

Execution lane: agent loop repeats until final answer; only search provider changes

100 hard questions
SealQA-Hard stratified sample

Provider condition
Brave / Tavily / Firecrawl
only varied component

Frozen GPT-5.4 agent
same prompt, tools, budget

Decision surface
document id, rank, title, URL,
snippet surface, metadata
no page text unless fetched

Agent loop step
observe evidence, reason,
then choose next action

Tool observations
search results or page evidence
returned to the agent

Final answer
short answer or abstention

fetch_page
selected document IDs
shared Jina Reader backend

search_web
one or many calls

Trajectory JSONL
all searches, visible URLs,
fetches, tokens, final + gold answer

Oracle lane: replay visible evidence one URL at a time

Judge records
one row per visible URL
page-visible rows include snippet fields

Kimi-K2.6 judge
gold support, contradiction,
model-answer support, spans,
garbage, confidence

Decision-surface metrics
pre-fetch / post-fetch support,
SMART/MISSED/BLIND/NO-OP,
surface c:g, semantic-correct